

Large-Scale Multimodal Datasets for Generalizable Plant Pathology: A Technical Comparative Analysis of Hybrid Controlled and Field Repositories

The transition of computer vision systems from highly controlled laboratory environments to unstructured, real-world agricultural fields represents one of the most significant challenges in modern smart farming¹. Early deep learning frameworks trained on idealized, single-leaf samples frequently exhibit catastrophic generalization decay when evaluated under realistic field conditions². The presence of complex, heterogeneous backgrounds, variable lighting conditions, soil and weed clutter, dynamic shadowing, and varying sensor specifications introduces a massive domain shift¹. To bridge this gap, the agricultural machine learning community has shifted toward curating massive, multi-crop datasets that intentionally combine clean laboratory images with noisy, in-the-wild field samples⁶. This technical report provides an exhaustive, peer-targeted analysis of the largest and most robust plant pathology datasets available, detailing their taxonomic composition, environmental distributions, visual-linguistic integration, and architectural implications for training highly generalizable diagnostic models⁹.

Technical Comparative Matrix of Large-Scale Hybrid Datasets

The following structured matrix compares the leading plant pathology datasets containing both controlled (laboratory) and in-the-wild (field) imagery, sorted by scale of original labeled images.

Dataset Name	Labeled Images	Distinct Crop Species	Fine-Grained Pathological Classes	Environmental Composition	Multimodal & Annotation Features	Public Accessibility & Provenance
LeafNet 2.0 (LeafMD)	278,131 total (255,855 core pairs)	37 crop species	197 crop-disease classes (144 disease	Controlled (Lab), Natural (In-the-wild), and Hybrid/Mi	Expert-curated captions, stage-specific annotation	Open access via Hugging Face under

			classes + healthy controls)	xed settings across 9 global regions	ns (early vs. late), symptom-focused text descriptions	CC-BY 4.0 International license; originated from UT Austin & Korea University ¹⁰
Deep-Plant-Disease	248,578	55 crop species	175 disease classes (333 unique crop-disease compositions)	Balanced controlled laboratory setups and real-world field conditions	Botanical taxonomy textual descriptions for crops and diseases, CLIP embeddings	Open access via Zenodo and GitHub; curated by PI@ntNet and CIRAD researchers ⁶
Plant Disease Dataset 271 (PDD271)	220,592	Multi-crop (Fruit Trees, Vegetables, Field Crops)	271 categories of plant pathologies	Combined laboratory backdrops and complex real-field environments	Visual patch clustering representations, hierarchical taxonomies	Restricted access due to corporate IP constraints (Beijing Puhui Sannong Technology Co. Ltd.); only sample subsets are open-sou

						rce ¹¹
LeafNet 1.0	186,000	22 crop species	97 fine-grained classes (62 disease categories + healthy controls)	Farm (In-the-wild) and Laboratory settings across 7 countries	Expert-annotated metadata, 13,950 question-answer pairs for Visual Question Answering (VQA)	Open access via Hugging Face and GitHub; predecessor of the LeafMD repository ¹⁴
FloraSyntropy Archive	178,922	35 crop species	97 distinct disease and healthy classes	Curation of 13 public repositories spanning 2018 to 2023 (controlled and wild mixes)	Optimized client node partitions for distributed federated learning configurations	Open access via PMC and GitHub under CC-BY 4.0; developed by DFKI researchers ¹⁶
PlantWild (v2 / MVPDR)	23,689 (v2) / 18,542 (v1)	34 plant hosts	115 categories (v2) / 89 categories (v1) of pathologies	Multi-viewpoint, variable-lighting, complex-background wild environments	Descriptive textual prompts from Wikipedia and GPT-3.5 verified by human pathologists	Open access via Google Drive, Hugging Face, and UQRDM; curated by University of Queensland


						nd research ers ¹
PlantVillage (Historical Baseline)	54,305 (or 54,303)	14 crop species	38 classes (17 fungal, 4 bacterial, 2 oomycete , 2 viral, 1 mite, 12 healthy)	Strictly controlled laborator y condition s (individua l leaves on uniform gray backdrop s)	Categoric al text labels only	Publicly archived across GitHub and Kaggle; original domain hosting is defunct ²⁰
PlantDoc	2,598 total (2,549 in cropped version)	13 plant species	17 to 27 classes dependin g on bounding box configurat ions	Scraped internet images captured under unstructu red field condition s	Bounding box coordinat es for object detection (8,595 labeled individual objects)	Open access via GitHub and Kaggle under CC-BY 4.0; develope d for benchma rking in-field models ²³

Detailed Curation and Structural Analysis of Prominent Repositories

LeafNet 2.0 (LeafMD)

LeafNet 2.0 represents the largest, most taxonomically complex multimodal agricultural dataset currently compiled¹⁰. Originating from research teams at the University of Texas at Austin and Korea University, LeafNet 2.0 was developed to address the spatial and environmental

limitations of predecessor models by significantly expanding geographic coverage to encompass nine distinct agricultural zones spanning tropical, subtropical, and temperate climates¹⁰.



The image dataset compiles 278,131 samples, from which a core of 255,855 standardized image-text pairs is derived¹⁰. Structurally, it integrates natural field environments, controlled laboratory settings, and hybrid/mixed backdrops captured using various devices such as smartphones and high-resolution digital cameras¹². Its core innovation lies in its multimodal pairing; rather than relying strictly on categorical disease labels, LeafNet 2.0 pairs each image with symptom-focused descriptions from plant pathology reference texts and disease stage annotations representing early and late-stage progression¹⁰. This design enables vision-language models (VLMs) to learn how visual symptoms morphologically evolve under heterogeneous field conditions¹². The evaluation portion, LeafBench 2.0, utilizes a label-constrained prompting strategy to test models across nine fine-grained pathological reasoning tasks, significantly minimizing open-ended generation hallucinations¹⁰.

Deep-Plant-Disease

Developed by researchers affiliated with the PI@ntNet consortium and CIRAD, the Deep-Plant-Disease dataset comprises 248,578 digital images across 55 crop species and 175 disease classes, yielding 333 unique crop-disease compositions⁶. The architectural goal of Deep-Plant-Disease is to foster zero-shot compositional generalization⁹. In real-world agriculture, collecting physical examples of every possible crop-pathogen pairing is practically impossible²⁹. Deep-Plant-Disease addresses this by enabling models to disentangle crop-specific features from disease-specific textures³⁰.

By packaging the visual database with standardized botanical taxonomy text descriptions for both the host plant and the associated pathogen, it provides a dual-branch visual-linguistic architecture with semantic anchors⁶. This allows contrastive vision-language models (such as CLIP-based variants) to accurately recognize unseen crop-disease pairs in the field by aligning visual pathology cues with text embedding vectors of the corresponding plant and disease⁹.

Plant Disease Dataset 271 (PDD271)

PDD271 is a highly structured, fine-grained plant leaf dataset containing 220,592 images mapping to 271 distinct pathological classes¹¹. The dataset is organized into three major macro-classes: Fruit Trees, Vegetables, and Field Crops¹¹.

Its environmental mix deliberately balances clean, tabletop laboratory views with highly complex natural backgrounds¹³. However, a significant operational limitation of PDD271 is its restrictive licensing¹¹. Belonging to Beijing Puhui Sannong Technology Co. Ltd., corporate intellectual property restrictions prevent the full dataset from being open-sourced¹¹. Only a small, representative sample subset is publicly downloadable, meaning that while it is highly cited for benchmark comparisons, it cannot serve as the primary training corpus for custom, large-scale open-source architectures¹¹.

FloraSyntropy Archive

The FloraSyntropy Archive compiles 178,922 images across 35 species and 97 disease/healthy classes, serving as a unified benchmark for evaluating model generalization¹⁶.

Recognizing that many existing agricultural networks achieve inflated accuracies on narrow, homogeneous databases, the creators of the FloraSyntropy Archive consolidated 13 distinct public repositories published between 2018 and 2023¹⁶. It integrates historical laboratory benchmarks, such as the original PlantVillage (specifically 20,639 selected samples) and Plant Village V2 (70,000 samples), with highly unstructured field data, such as Cassava (53,303 samples) and localized datasets like BananaLSD¹⁶.

The primary design objective of FloraSyntropy is to evaluate privacy-preserving, scalable federated learning frameworks³⁴. The archive's metadata partitions are engineered to simulate local client nodes (such as regional farms or distinct testing stations), evaluating how client models trained on localized edge data generalize when aggregated via a central server using federated optimization algorithms¹⁷.

Technical Implications of Controlled and Wild Coexistence

The strategic integration of both laboratory and in-the-wild images within a single training pipeline represents a core conceptual advancement in agricultural AI³⁶. Understanding the causal relationships and underlying feature dynamics of this environmental mix explains why this approach is highly effective.

Feature Disentanglement and Attentional Shift

When models are trained exclusively on controlled laboratory images, such as the historical

PlantVillage baseline, the visual feature space is highly simplified⁴. Leaves are flat, centered, and isolated against uniform, high-contrast backgrounds⁴. Consequently, deep neural network optimization paths easily align with superficial global visual shortcuts, such as the overall shape of the leaf margin or background-color borders, rather than learning the actual visual pathology of the lesions¹³. When these models are deployed under real field conditions, their classification performance collapses³.

Conversely, training exclusively on small, internet-scraped in-the-wild datasets (such as PlantDoc) often results in poor convergence and high variance¹. The abundance of visual noise—such as soil patches, dynamic sun glints, shadows, dew drops, insect damage, and branch occlusions—can obscure the fine-grained, localized visual cues of early-stage fungal pustules or bacterial chlorosis⁵.

Combining controlled laboratory images and natural field images in massive repositories establishes a structured representation learning process:

- **Laboratory Images as High-Fidelity Representation Anchors:** Clean, controlled-environment images allow the visual encoder's feature space to map the pure, unoccluded visual syntax of standard disease lesions (such as necrotic spots, chlorotic halos, and powdery mildew textures) without background interference⁴.
- **In-the-Wild Images as Domain Regularizers:** Complex-background field images act as an adversarial regularization mechanism³². They force the model to discard global tabletop spatial heuristics and learn domain-invariant, localized lesion patterns regardless of environmental noise, lighting shifts, or sensor degradation³⁰.

Through explainability techniques like Gradient-weighted Class Activation Mapping (Grad-CAM), agricultural AI researchers have verified that this coexistence fundamentally alters model attention³⁰. While pure convolutional neural networks (CNNs) tend to focus heavily on broad, structural plant contours, the integration of vision transformers (ViTs) and hybrid attention networks (such as PlantAIM) allows the model to simultaneously process global leaf morphology and zoom in on localized disease lesions³⁰.

Dataset Discrepancies and Provenance Challenges

Practitioners must maintain rigorous quality control when downloading public datasets, as community-driven repositories often introduce structural discrepancies³⁶. A notable example exists within the popular *PlantWild* repository¹.

The original, peer-reviewed PlantWild dataset published at ACM Multimedia 2024 is documented as an in-the-wild, multimodal plant disease dataset curated from diverse internet search engines (Google, Bing, Baidu) under Creative Commons licenses, containing 18,542 real-world images across 89 classes¹.


However, community-hosted variations on public model hubs (such as Hugging Face and personal GitHub repos) often repackage these files with altered annotations³⁶. In some open-source project implementations, the "PlantWild" subset is incorrectly labeled as a controlled laboratory database with uniform backgrounds (18,533 images), while being paired with "FieldPlant" (5,156 images) to represent the wild field domain³⁶.

This structural variance highlights the absolute necessity of verifying dataset checksums, licensing agreements, and official publication metadata before initiating large-scale training pipelines to ensure reproducibility and legal compliance⁷.

Semantic Standardization via EPPO Taxonomy Integration

A persistent operational bottleneck in aggregating multi-source agricultural datasets is the fragmentation of pathological nomenclature⁷. Across global databases, the same pathogen or disease is frequently annotated under different common names, synonyms, or distinct life cycle stages⁷. For example, a single fungal infection might be cataloged as "Apple Scab," "Venturia inaequalis" (its taxonomic binomial name), or "scab-frog-eye complex" depending on whether it was sourced from PlantVillage or the Plant Pathology 2021 challenge²¹.

To enable models to generalize across global agricultural domains, researchers have developed automated pipelines that utilize large language models (LLMs) such as LLaMa and ChatGPT to standardize all dataset classifications⁷.



By querying the European and Mediterranean Plant Protection Organization (EPPO) Global Database, these AI pipelines resolve conflicting taxonomic synonyms and establish a standardized binomial nomenclature⁷. This taxonomy integration yields significant operational advantages:

- **Interoperable Metadata Concurrence:** It allows researchers to seamlessly combine disparate, localized datasets into massive, highly structured repositories (like LeafNet 2.0) without creating redundant or overlapping output classes⁷.
- **Hierarchical Taxonomic Reasoning:** Standardizing labels according to biological

hierarchy (Kingdom → Phylum → Class → Order → Family → Genus →

Species) allows vision-language models to perform hierarchical inferences¹⁰. If a model cannot definitively identify a rare, unseen pathogen species, it can still leverage the underlying semantic text embeddings to predict the correct genus or family, avoiding catastrophic classification failure in the field¹⁰.

Downstream Modeling and Implementation Roadmap

Selecting the optimal dataset and modeling paradigm depends directly on the hardware constraints, training budgets, and deployment targets of the agricultural system.

Option A: Training Large-Scale Vision-Language and Diagnostic Reasoning Models

For developers building interactive, chat-enabled diagnostics or zero-shot decision support systems, **LeafNet 2.0 (LeafMD)** is the ideal foundation dataset¹⁰.

- **Implementation Strategy:** The model should be trained using multimodal vision-language architectures, aligning visual encoders (such as ViT-L or SigLIP) with autoregressive language models (such as PaliGemma or LLaVA)¹⁰.
- **Algorithmic Focus:** Utilizing the label-constrained multiple-choice QA pairs in LeafBench 2.0 prevents free-form text hallucination¹⁴. The model must be optimized on tasks like Pathogen Classification (PC) and Disease Severity Classification (DSC) to ensure it performs fine-grained clinical reasoning rather than relying on coarse visual correlations¹⁰.

Option B: Zero-Shot Generalization to Unseen Crops and Emerging Diseases

For applications requiring robust zero-shot generalization to foreign agricultural regions or newly mutated crop diseases, **Deep-Plant-Disease** represents the optimal choice⁶.

- **Implementation Strategy:** Developers should implement contrastive architectures like **CL-ViT** (Cross-Learning Vision Transformer) or **FF-CLIP**⁹.
- **Algorithmic Focus:** The model leverages a dual-branch visual encoder to separately capture disease-specific semantic representations and crop-domain representations³¹. It utilizes a frozen CLIP text encoder to guide visual feature extraction through text-based semantic anchoring, enabling the network to identify known disease patterns on crop species that were never present in the training set³¹.

Option C: Efficient Edge and Offline Mobile Deployment

For models deployed directly on smartphones, edge sensors, or agricultural drone payloads where cloud connectivity is unavailable or bandwidth-constrained, a combination of **PlantWild v2** and **FloraSyntropy Archive** is recommended¹⁹.

- **Implementation Strategy:** The framework should leverage highly compressed, lightweight architectures (such as MobileNetV3-Small, ConvNeXt-Tiny, or YOLOv11x)³⁶.
- **Algorithmic Focus:** Rather than training compact models from scratch on noisy field data, developers should use **Knowledge Distillation**⁴³. In this paradigm, a high-capacity, heavy vision transformer (the teacher, pre-trained on a massive database like

PI@ntNet-300K or iNaturalist 2021) transfers its rich, domain-invariant representation space to the lightweight student model⁴³. This allows the compact edge network to retain the fine-grained discriminative capabilities of large models while maintaining low-latency and energy-efficient inference in the field⁴³.

Technical Conclusions

Developing automated computer vision tools capable of reliably identifying foliar pathologies in diverse environmental settings requires datasets that reflect the actual complexity of global agriculture¹⁶. The traditional reliance on isolated, single-domain databases is being replaced by comprehensive, multimodal repositories that bridge the gap between idealized laboratory controls and unstructured, real-world fields⁸.

By compiling hundreds of thousands of images across diverse geographic regions, and standardizing their nomenclatures via LLM-driven taxonomic pipelines, datasets like **LeafNet 2.0** and **Deep-Plant-Disease** provide the foundational data scale and representation depth required to train highly generalizable, robust, and field-ready agricultural models⁷.

Works cited

1. Benchmarking In-the-Wild Multimodal Plant Disease Recognition and A Versatile Baseline - arXiv, <https://arxiv.org/html/2408.03120v1>
2. Real-Field-Ready and Digitally Sustainable Plant Disease Recognition via Federated Multimodal Edge Learning and Few-Shot Domain Adaptation - MDPI, <https://www.mdpi.com/2077-0472/16/9/918>
3. DIANA: A deep learning-based paprika plant disease and pest phenotyping system with disease severity analysis - Frontiers, <https://www.frontiersin.org/journals/plant-science/articles/10.3389/fpls.2022.983625/full>
4. FieldPlant Dataset for Plant Disease Detection | PDF | Deep Learning | Maize - Scribd, <https://www.scribd.com/document/696389577/FieldPlant-A-dataset-of-field-plant-images-for-pla>
5. A general-purpose edge-feature guidance module to enhance vision transformers for plant disease identification | Request PDF - ResearchGate, https://www.researchgate.net/publication/374092310_A_general-purpose_edge-feature_guidance_module_to_enhance_vision_transformers_for_plant_disease_identification
6. abelchai/Deep-Plant-Disease-Dataset-Is-All-You-Need-for-Plant-Disease-Identification - GitHub, <https://github.com/abelchai/Deep-Plant-Disease-Dataset-Is-All-You-Need-for-Plant-Disease-Identification>
7. Identifying plant diseases using artificial intelligence - Cirad, <https://www.cirad.fr/en/cirad-news/news/2026/identifying-plant-diseases>
8. Visual Information Guided Multi-modal Model for Plant Disease Anomaly Detection,

- https://www.researchgate.net/publication/384170118_Visual_Information_Guided_Multi-modal_Model_for_Plant_Disease_Anomaly_Detection
9. GitHub - abelchai/Cross-Learning-Vision-Transformer-CL-ViT, <https://github.com/abelchai/Cross-Learning-Vision-Transformer-CL-ViT>
 10. A multiregional image-text dataset and benchmark for vision-language modeling of plant diseases | bioRxiv, <https://www.biorxiv.org/content/10.64898/2026.07.01.735881v2>
 11. liuxindazz/PDD271 - GitHub, <https://github.com/liuxindazz/PDD271>
 12. A multiregional image-text dataset and benchmark for vision- language modeling of plant diseases - bioRxiv, <https://www.biorxiv.org/content/biorxiv/early/2026/07/02/2026.07.01.735881.full.pdf>
 13. Plant Disease Recognition: A Large-Scale Benchmark Dataset and a Visual Region and Loss Reweighting Approach - ijrti, <https://www.ijrti.org/papers/IJRTI2209006.pdf>
 14. LeafNet: A Large-Scale Dataset and Comprehensive Benchmark for Foundational Vision-Language Understanding of Plant Diseases - arXiv, <https://arxiv.org/html/2602.13662v2>
 15. enalis/LeafBench · Datasets at Hugging Face, <https://huggingface.co/datasets/enalis/LeafBench>
 16. Scalable Deep Learning with Novel FloraSyntropy Archive for Large-Scale Plant Disease Diagnosis - arXiv, <https://arxiv.org/html/2508.17653v1>
 17. FloraSyntropy-net: scalable deep learning with novel FloraSyntropy archive for large-scale plant disease diagnosis - PMC, <https://pmc.ncbi.nlm.nih.gov/articles/PMC13064331/>
 18. MVPDR/README.md at main - GitHub, <https://github.com/tgwei05/MVPDR/blob/main/README.md>
 19. uqtwei2/PlantWild · Datasets at Hugging Face, <https://huggingface.co/datasets/uqtwei2/PlantWild>
 20. Android Application for Plant Disease Detection Using this Deep Learning - International Journal of Advanced Research in Science, Communication and Technology (IJARSCT), <https://ijarsct.co.in/Paper27657.pdf>
 21. Leaf-Based Plant Disease Detection and Explainable AI - arXiv, <https://arxiv.org/html/2404.16833v1>
 22. gabrieldgf4/PlantVillage-Dataset - GitHub, <https://github.com/gabrieldgf4/PlantVillage-Dataset>
 23. Plant Doc Dataset - Kaggle, <https://www.kaggle.com/datasets/abdulhasibuddin/plant-doc-dataset>
 24. PlantDoc - Dataset Ninja, <https://datasetninja.com/plantdoc>
 25. PlantDoc: A Dataset for Visual Plant Disease Detection - GitHub, <https://github.com/pratikkayal/PlantDoc-Dataset>
 26. A multiregional image-text dataset and benchmark for vision-language modeling of plant diseases - ResearchGate, https://www.researchgate.net/publication/408383990_A_multiregional_image-text_dataset_and_benchmark_for_vision-language_modeling_of_plant_diseases

27. enalis/LeafBenchV2 · Datasets at Hugging Face,
<https://huggingface.co/datasets/enalis/LeafBenchV2>
28. A multiregional image-text dataset and benchmark for vision-language modeling of plant diseases | bioRxiv,
<https://www.biorxiv.org/content/10.64898/2026.07.01.735881v1>
29. abelchai/FF-CLIP-Can-Language-Improve-Visual-Features-For-Distinguishing-Unseen-Plant-Diseases - GitHub,
<https://github.com/abelchai/FF-CLIP-Can-Language-Improve-Visual-Features-For-Distinguishing-Unseen-Plant-Diseases>
30. abelchai/PlantAIM: Plant Disease Global-Local Features Fusion Attention Model - GitHub, <https://github.com/abelchai/PlantAIM>
31. Pairwise Feature Learning for Unseen Plant Disease Recognition - ResearchGate,
https://www.researchgate.net/publication/374543673_Pairwise_Feature_Learning_for_Unseen_Plant_Disease_Recognition
32. Plant Disease Recognition: A Large-Scale Benchmark Dataset and a Visual Region and Loss Reweighting Approach,
<http://vipl.ict.ac.cn/homepage/sqjiang/Publication/Journal/202206/P020220602347207892237.pdf>
33. Plant Disease Recognition: A Visual Region and Loss Reweighting Approach - ijarst, <https://ijarst.co.in/Paper9035.pdf>
34. Development of Deep Learning Technique for Crop Pest Classification | Request PDF,
https://www.researchgate.net/publication/394210648_Development_of_Deep_Learning_Technique_for_Crop_Pest_Classification
35. Illustration of the proposed federated learning approach. Farms... - ResearchGate,
https://www.researchgate.net/figure/Illustration-of-the-proposed-federated-learning-approach-Farms-optional-aggregate-a_fig4_391167628
36. JK-TK/PlantDiseaseDetection - Hugging Face,
<https://huggingface.co/JK-TK/PlantDiseaseDetection>
37. PlantVillage - Kaggle,
<https://www.kaggle.com/datasets/mohitsingh1804/plantvillage>
38. DIANA: A deep learning-based paprika plant disease and pest phenotyping system with disease severity analysis - PMC,
<https://pmc.ncbi.nlm.nih.gov/articles/PMC9582859/>
39. Snap and Diagnose: An Advanced Multimodal Retrieval System for Identifying Plant Diseases in the Wild - arXiv, <https://arxiv.org/html/2408.14723v1>
40. A Large-Scale In-the-wild Dataset for Plant Disease Segmentation - PMC,
<https://pmc.ncbi.nlm.nih.gov/articles/PMC12891534/>
41. [2602.13662] LeafNet: A Large-Scale Dataset and Comprehensive Benchmark for Foundational Vision-Language Understanding of Plant Diseases - arXiv,
<https://arxiv.org/abs/2602.13662>
42. Can Language Improve Visual Features For Distinguishing Unseen Plant Diseases? | Request PDF - ResearchGate,
https://www.researchgate.net/publication/386421969_Can_Language_Improve_Visual_Features_For_Distinguishing_Unseen_Plant_Diseases

43. Energy-Efficient Plant Monitoring via Knowledge Distillation - arXiv,
<https://arxiv.org/html/2604.27178v2>
44. (PDF) Towards Sustainable Agricultural Systems: A Lightweight Deep Learning Model for Plant Disease Detection - ResearchGate,
https://www.researchgate.net/publication/371381162_Towards_Sustainable_Agricultural_Systems_A_Lightweight_Deep_Learning_Model_for_Plant_Disease_Detection
45. Energy-Efficient Plant Monitoring via Knowledge Distillation - arXiv,
<https://arxiv.org/pdf/2604.27178>
46. Self-Supervised Learning of Plant Image Representations - arXiv,
<https://arxiv.org/pdf/2604.27538>